

Captcha_FixCaptcha

April 20, 2025

```
[1]: import os
import glob
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms
from torchvision.models import resnet18

import string
from tqdm.notebook import tqdm
import cv2
from PIL import Image
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import multiprocessing as mp
```

```
[2]: cpu_count = mp.cpu_count()
print(cpu_count)
```

2

```
[7]: data_path = "/content/drive/MyDrive/archive (1)/samples"
```

```
[9]: image_fns = os.listdir(data_path)
print(len(image_fns))
print(np.unique([len(image_fn.split(".")[0]) for image_fn in image_fns]))
```

1080

[5 9]

```
[10]: for idx, image_fn in enumerate(image_fns):
    if len(image_fn.split(".")[0]) != 5:
        print(idx, image_fn)
```

```

2 3mxdn (1).png
3 3fbxd (1).png
5 3n7mx (1).png
7 3n3cf (1).png
10 3eny7 (1).png
17 3n2b4 (1).png
20 3ndxd (1).png
1076 3g2w6 (1).png
1077 3ebnn (1).png
1078 3ebpw (1).png

```

```
[12]: image_fns_train, image_fns_test = train_test_split(image_fns, random_state=0)
      print(len(image_fns_train), len(image_fns_test))
```

```
810 270
```

```
[13]: image_ns = [image_fn.split(".")[0] for image_fn in image_fns]
      image_ns = "".join(image_ns)
      letters = sorted(list(set(list(image_ns))))
      print(len(letters))
      print(letters)
```

```
23
```

```
[' ', '(', ')', '1', '2', '3', '4', '5', '6', '7', '8', 'b', 'c', 'd', 'e', 'f',
'g', 'm', 'n', 'p', 'w', 'x', 'y']
```

```
[14]: vocabulary = ["-"] + letters
      print(len(vocabulary))
      print(vocabulary)
      idx2char = {k:v for k,v in enumerate(vocabulary, start=0)}
      print(idx2char)
      char2idx = {v:k for k,v in idx2char.items()}
      print(char2idx)
```

```
24
```

```
['-', ' ', '(', ')', '1', '2', '3', '4', '5', '6', '7', '8', 'b', 'c', 'd', 'e',
'f', 'g', 'm', 'n', 'p', 'w', 'x', 'y']
{0: '-', 1: ' ', 2: '(', 3: ')', 4: '1', 5: '2', 6: '3', 7: '4', 8: '5', 9: '6',
10: '7', 11: '8', 12: 'b', 13: 'c', 14: 'd', 15: 'e', 16: 'f', 17: 'g', 18: 'm',
19: 'n', 20: 'p', 21: 'w', 22: 'x', 23: 'y'}
{'-': 0, ' ': 1, '(': 2, ')': 3, '1': 4, '2': 5, '3': 6, '4': 7, '5': 8, '6': 9,
'7': 10, '8': 11, 'b': 12, 'c': 13, 'd': 14, 'e': 15, 'f': 16, 'g': 17, 'm': 18,
'n': 19, 'p': 20, 'w': 21, 'x': 22, 'y': 23}
```

```
[15]: batch_size = 16
```

```
[16]: class CAPTCHADataset(Dataset):
      def __init__(self, data_dir, image_fns):
```

```

        self.data_dir = data_dir
        self.image_fns = image_fns

    def __len__(self):
        return len(self.image_fns)

    def __getitem__(self, index):
        image_fn = self.image_fns[index]
        image_fp = os.path.join(self.data_dir, image_fn)
        image = Image.open(image_fp).convert('RGB')
        image = self.transform(image)
        text = image_fn.split(".")[0]
        return image, text

    def transform(self, image):

        transform_ops = transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize(mean=(0.485, 0.456, 0.406), std=(0.229, 0.224, ↵
↵0.225))
        ])
        return transform_ops(image)

```

```

[17]: trainset = CAPTCHADataset(data_path, image_fns_train)
      testset = CAPTCHADataset(data_path, image_fns_test)
      train_loader = DataLoader(trainset, batch_size=batch_size, ↵
↵num_workers=cpu_count, shuffle=True)
      test_loader = DataLoader(testset, batch_size=batch_size, num_workers=cpu_count, ↵
↵shuffle=False)
      print(len(train_loader), len(test_loader))

```

51 17

```

[21]: image_batch, text_batch = next(iter(train_loader))
      print(image_batch.size(), text_batch)

```

```

torch.Size([16, 3, 50, 200]) ('78dw6', '3n2b4 (1)', 'pcpg6', '56m6y', '4fp5g',
'byfgn', 'g78gn', 'mdxpn', 'bpwd7', '8n34n', 'f2fge', 'b4ndb', 'nwg2m', 'gp7c5',
'pg2yx', '3n7mx')

```

```

[22]: num_chars = len(char2idx)
      print(num_chars)
      rnn_hidden_size = 256

```

24

```

[23]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
      print(device)

```

cpu

```
[24]: resnet = resnet18(pretrained=True)
```

```
/usr/local/lib/python3.11/dist-packages/torchvision/models/_utils.py:208:
UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/torchvision/models/_utils.py:223:
UserWarning: Arguments other than a weight enum or `None` for 'weights' are
deprecated since 0.13 and may be removed in the future. The current behavior is
equivalent to passing `weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to
/root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100%|          | 44.7M/44.7M [00:00<00:00, 96.1MB/s]
```

```
[25]: print(resnet)
```

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
```

```

track_running_stats=True)
    )
    )
    (layer2): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (1): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (layer3): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    )
  )
)

```

```

    )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (layer4): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc): Linear(in_features=512, out_features=1000, bias=True)
)

```

```
[26]: class CRNN(nn.Module):

    def __init__(self, num_chars, rnn_hidden_size=256, dropout=0.1):

        super(CRNN, self).__init__()
        self.num_chars = num_chars
        self.rnn_hidden_size = rnn_hidden_size
        self.dropout = dropout

        # CNN Part 1
        resnet_modules = list(resnet.children())[:-3]
        self.cnn_p1 = nn.Sequential(*resnet_modules)

        # CNN Part 2
        self.cnn_p2 = nn.Sequential(
            nn.Conv2d(256, 256, kernel_size=(3,6), stride=1, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True)
        )
        self.linear1 = nn.Linear(1024, 256)

        # RNN
        self.rnn1 = nn.GRU(input_size=rnn_hidden_size,
                           hidden_size=rnn_hidden_size,
                           bidirectional=True,
                           batch_first=True)
        self.rnn2 = nn.GRU(input_size=rnn_hidden_size,
                           hidden_size=rnn_hidden_size,
                           bidirectional=True,
                           batch_first=True)
        self.linear2 = nn.Linear(self.rnn_hidden_size*2, num_chars)

    def forward(self, batch):

        batch = self.cnn_p1(batch)
        # print(batch.size()) # torch.Size([-1, 256, 4, 13])

        batch = self.cnn_p2(batch) # [batch_size, channels, height, width]
        # print(batch.size())# torch.Size([-1, 256, 4, 10])

        batch = batch.permute(0, 3, 1, 2) # [batch_size, width, channels, height]
        # print(batch.size()) # torch.Size([-1, 10, 256, 4])

        batch_size = batch.size(0)
        T = batch.size(1)
```

```

        batch = batch.view(batch_size, T, -1) # [batch_size, T==width,
↪num_features==channels*height]
        # print(batch.size()) # torch.Size([-1, 10, 1024])

        batch = self.linear1(batch)
        # print(batch.size()) # torch.Size([-1, 10, 256])

        batch, hidden = self.rnn1(batch)
        feature_size = batch.size(2)
        batch = batch[:, :, :feature_size//2] + batch[:, :, feature_size//2:]
        # print(batch.size()) # torch.Size([-1, 10, 256])

        batch, hidden = self.rnn2(batch)
        # print(batch.size()) # torch.Size([-1, 10, 512])

        batch = self.linear2(batch)
        # print(batch.size()) # torch.Size([-1, 10, 20])

        batch = batch.permute(1, 0, 2) # [T==10, batch_size,
↪num_classes==num_features]
        # print(batch.size()) # torch.Size([10, -1, 20])

    return batch

```

```

[27]: def weights_init(m):
        classname = m.__class__.__name__
        if type(m) in [nn.Linear, nn.Conv2d, nn.Conv1d]:
            torch.nn.init.xavier_uniform_(m.weight)
            if m.bias is not None:
                m.bias.data.fill_(0.01)
        elif classname.find('BatchNorm') != -1:
            m.weight.data.normal_(1.0, 0.02)
            m.bias.data.fill_(0)

```

```

[28]: crnn = CRNN(num_chars, rnn_hidden_size=rnn_hidden_size)
        crnn.apply(weights_init)
        crnn = crnn.to(device)

```

```

[29]: text_batch_logits = crnn(image_batch.to(device))
        print(text_batch)
        print(text_batch_logits.shape)

```

```

('78dw6', '3n2b4 (1)', 'pcpg6', '56m6y', '4fp5g', 'byfgn', 'g78gn', 'mdxpn',
'bpwd7', '8n34n', 'f2fge', 'b4ndb', 'nwg2m', 'gp7c5', 'pg2yx', '3n7mx')
torch.Size([10, 16, 24])

```

```

[30]: criterion = nn.CTCLoss(blank=0)

```



```
[31]: def encode_text_batch(text_batch):
```

```
    text_batch_targets_lens = [len(text) for text in text_batch]
    text_batch_targets_lens = torch.IntTensor(text_batch_targets_lens)

    text_batch_concat = "".join(text_batch)
    text_batch_targets = [char2idx[c] for c in text_batch_concat]
    text_batch_targets = torch.IntTensor(text_batch_targets)

    return text_batch_targets, text_batch_targets_lens
```

```
[32]: def compute_loss(text_batch, text_batch_logits):
```

```
    """
    text_batch: list of strings of length equal to batch size
    text_batch_logits: Tensor of size([T, batch_size, num_classes])
    """

    text_batch_logits = F.log_softmax(text_batch_logits, 2) # [T, batch_size, ↵
    ↵num_classes]
    text_batch_logits_lens = torch.full(size=(text_batch_logits.size(1),),
                                       fill_value=text_batch_logits.size(0),
                                       dtype=torch.int32).to(device) # ↵
    ↵[batch_size]
    #print(text_batch_logits.shape)
    #print(text_batch_logits_lens)
    text_batch_targets, text_batch_targets_lens = encode_text_batch(text_batch)
    #print(text_batch_targets)
    #print(text_batch_targets_lens)
    loss = criterion(text_batch_logits, text_batch_targets, ↵
    ↵text_batch_logits_lens, text_batch_targets_lens)

    return loss
```

```
[33]: compute_loss(text_batch, text_batch_logits)
```

```
[33]: tensor(4.6650, grad_fn=<MeanBackward0>)
```

```
[35]: num_epochs = 50
```

```
lr = 0.001
```

```
weight_decay = 1e-3
```

```
clip_norm = 5
```

```
[37]: optimizer = optim.Adam(crn.parameters(), lr=lr, weight_decay=weight_decay)
```

```
lr_scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, verbose=True, ↵
    ↵patience=5)
```

```
[38]: crnn = CRNN(num_chars, rnn_hidden_size=rnn_hidden_size)
```

```
crnn.apply(weights_init)
```

```
crnn = crnn.to(device)
```

```
[39]: epoch_losses = []
iteration_losses = []
num_updates_epochs = []
for epoch in tqdm(range(1, num_epochs+1)):
    epoch_loss_list = []
    num_updates_epoch = 0
    for image_batch, text_batch in tqdm(train_loader, leave=False):
        optimizer.zero_grad()
        text_batch_logits = crnn(image_batch.to(device))
        loss = compute_loss(text_batch, text_batch_logits)
        iteration_loss = loss.item()

        if np.isnan(iteration_loss) or np.isinf(iteration_loss):
            continue

        num_updates_epoch += 1
        iteration_losses.append(iteration_loss)
        epoch_loss_list.append(iteration_loss)
        loss.backward()
        nn.utils.clip_grad_norm_(crnn.parameters(), clip_norm)
        optimizer.step()

    epoch_loss = np.mean(epoch_loss_list)
    print("Epoch:{} Loss:{} NumUpdates:{}".format(epoch, epoch_loss,
↪num_updates_epoch))
    epoch_losses.append(epoch_loss)
    num_updates_epochs.append(num_updates_epoch)
    lr_scheduler.step(epoch_loss)
```

```
0%|          | 0/50 [00:00<?, ?it/s]
```

```
0%|          | 0/51 [00:00<?, ?it/s]
```

```
Epoch:1    Loss:3.0916408090030445    NumUpdates:51
```

```
0%|          | 0/51 [00:00<?, ?it/s]
```

```
Epoch:2    Loss:2.4448012838176654    NumUpdates:51
```

```
0%|          | 0/51 [00:00<?, ?it/s]
```

```
Epoch:3    Loss:1.9191021241393744    NumUpdates:51
```

```
0%|          | 0/51 [00:00<?, ?it/s]
```

```
Epoch:4    Loss:1.6287312788121842    NumUpdates:51
```

```
0%|          | 0/51 [00:00<?, ?it/s]
```

```
Epoch:5    Loss:1.462795521698746    NumUpdates:51
```

0%	0/51 [00:00<?, ?it/s]	
Epoch:6	Loss:1.338064009068059	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:7	Loss:1.2598869286331476	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:8	Loss:1.1965055465698242	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:9	Loss:1.1697289639828252	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:10	Loss:1.134702811054155	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:11	Loss:1.13827491274067	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:12	Loss:1.074782133102417	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:13	Loss:1.047883411248525	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:14	Loss:1.0220145910393958	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:15	Loss:0.9577525573618272	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:16	Loss:0.8911933723618003	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:17	Loss:0.8617152314560086	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:18	Loss:0.8544135900104747	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:19	Loss:0.8854422499151791	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:20	Loss:0.966071111314437	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:21	Loss:0.8648616204074785	NumUpdates:51

0%	0/51 [00:00<?, ?it/s]	
Epoch:22	Loss:0.8023002206110487	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:23	Loss:0.7963328618629306	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:24	Loss:0.8482310070711023	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:25	Loss:0.8287536209704829	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:26	Loss:0.8210709305370555	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:27	Loss:0.7841405611412198	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:28	Loss:0.7642733326145247	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:29	Loss:0.7630844875877979	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:30	Loss:0.7558378530483619	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:31	Loss:0.7394053316583821	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:32	Loss:0.7397293647130331	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:33	Loss:0.7359461270126642	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:34	Loss:0.7334616149173063	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:35	Loss:0.912491943321976	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:36	Loss:0.857014349862641	NumUpdates:51
0%	0/51 [00:00<?, ?it/s]	
Epoch:37	Loss:0.7864211680842381	NumUpdates:51

```

0%|          | 0/51 [00:00<?, ?it/s]
Epoch:38    Loss:0.7335886768266267    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:39    Loss:0.708309413171282    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:40    Loss:0.7032125019559673    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:41    Loss:0.6993528288953444    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:42    Loss:0.6973183973162782    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:43    Loss:0.7001666940894782    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:44    Loss:0.6976179211747413    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:45    Loss:0.698658942007551    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:46    Loss:0.7139831839823255    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:47    Loss:0.7137583131883659    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:48    Loss:0.8418545652838314    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:49    Loss:0.7703956272087845    NumUpdates:51
0%|          | 0/51 [00:00<?, ?it/s]
Epoch:50    Loss:0.6744280749676275    NumUpdates:51

```

```
[40]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 5))
```

```


ax1.plot(epoch_losses)
ax1.set_xlabel("Epochs")
ax1.set_ylabel("Loss")

ax2.plot(iteration_losses)
ax2.set_xlabel("Iterations")

```

```
ax2.set_ylabel("Loss")

plt.show()
```



```
[41]: def decode_predictions(text_batch_logits):

    text_batch_tokens = F.softmax(text_batch_logits, 2).argmax(2) # [T, ↵
    ↪batch_size]
    text_batch_tokens = text_batch_tokens.numpy().T # [batch_size, T]

    text_batch_tokens_new = []
    for text_tokens in text_batch_tokens:
        text = [idx2char[idx] for idx in text_tokens]
        text = "".join(text)
        text_batch_tokens_new.append(text)

    return text_batch_tokens_new

[42]: results_train = pd.DataFrame(columns=['actual', 'prediction'])
train_loader = DataLoader(trainset, batch_size=16, num_workers=1, shuffle=False)
with torch.no_grad():
    for image_batch, text_batch in tqdm(train_loader, leave=True):
        text_batch_logits = crnn(image_batch.to(device)) # [T, batch_size, ↵
        ↪num_classes==num_features]
        text_batch_pred = decode_predictions(text_batch_logits.cpu())
        #print(text_batch, text_batch_pred)
        df = pd.DataFrame(columns=['actual', 'prediction'])
```

```

        df['actual'] = text_batch
        df['prediction'] = text_batch_pred
        results_train = pd.concat([results_train, df])
results_train = results_train.reset_index(drop=True)

```

```
0%|          | 0/51 [00:00<?, ?it/s]
```

```

[43]: results_test = pd.DataFrame(columns=['actual', 'prediction'])
test_loader = DataLoader(testset, batch_size=16, num_workers=1, shuffle=False)
with torch.no_grad():
    for image_batch, text_batch in tqdm(test_loader, leave=True):
        text_batch_logits = crnn(image_batch.to(device)) # [T, batch_size, num_classes==num_features]
        text_batch_pred = decode_predictions(text_batch_logits.cpu())
        #print(text_batch, text_batch_pred)
        df = pd.DataFrame(columns=['actual', 'prediction'])
        df['actual'] = text_batch
        df['prediction'] = text_batch_pred
        results_test = pd.concat([results_test, df])
results_test = results_test.reset_index(drop=True)

```

```
0%|          | 0/17 [00:00<?, ?it/s]
```

```

[44]: print(results_train.shape)
results_train.head()

```

```
(810, 2)
```

```

[44]:  actual  prediction
0  ndecc  nnddeccccc
1  gecmf  ggeeccmmff
2  325fb  332255ffbb
3  6f857  66ff885577
4  87nym  8877nnyymm

```

```

[45]: print(results_test.shape)
results_test.head()

```

```
(270, 2)
```

```

[45]:  actual  prediction
0  x362g  xx336622gg
1  xemyg  xxeemmyygg
2  mye68  mmyyee6688
3  fbp2c  ffbppp22cc
4  mggce  mmggggccee

```

```

[46]: def remove_duplicates(text):
        if len(text) > 1:

```

```

        letters = [text[0]] + [letter for idx, letter in enumerate(text[1:],
↪start=1) if text[idx] != text[idx-1]]
        elif len(text) == 1:
            letters = [text[0]]
        else:
            return ""
        return "".join(letters)

def correct_prediction(word):
    parts = word.split("-")
    parts = [remove_duplicates(part) for part in parts]
    corrected_word = "".join(parts)
    return corrected_word

```

```

[47]: results_train['prediction_corrected'] = results_train['prediction'].
↪apply(correct_prediction)
results_train.head()

```

```

[47]:  actual  prediction prediction_corrected
0  ndecc  nnddeecccc          ndec
1  gecmf  ggeecmmmff          gecmf
2  325fb  332255ffbb          325fb
3  6f857  66ff885577          6f857
4  87nym  8877nnyymm          87nym

```

```

[48]: results_test['prediction_corrected'] = results_test['prediction'].
↪apply(correct_prediction)
results_test.head()

```

```

[48]:  actual  prediction prediction_corrected
0  x362g  xx336622gg          x362g
1  xemyg  xxeemmyygg          xemyg
2  mye68  mmyyee6688          mye68
3  fbp2c  ffbbpp22cc          fbp2c
4  mggce  mmggggccee          mgce

```

```

[49]: mistakes_df = results_test[results_test['actual'] !=
↪results_test['prediction_corrected']]
mistakes_df

```

```

[49]:  actual  prediction prediction_corrected
4      mggce  mmggggccee          mgce
5      y5dpp  yy55ddpppp          y5dp
6      3nfdn  33nnndddnn          3ndn
9      e2d66  ee22dd6666          e2d6
19     enn7n  eennnn77nn          en7n
20     x44n4  xx4444nn44          x4n4

```


30	77387	7777788877	787
47	8cccc	88c-----cc	8cc
49	n3ffn	nn33ffffnn	n3fn
53	8bbm4	88bbbbmm44	8bm4
62	cdcb3	eeddccb33	edcb3
63	d75b5	dd77555555	d75
73	ppwyd	ppppwwyydd	pwyd
90	mmc5n	mmmmmmccnn	mcn
94	m3588	mm33558888	m358
102	gfxcc	ggffxxcccc	gxfc
110	yyn57	yyyynn5577	yn57
111	474ff	447744ffff	474f
136	p8wvf	pp88wwwvff	p8wf
144	3n7mx (1)	33nn77mmxx	3n7mx
146	mmfm6	mmmmmmffmm	mfm
147	3mxdn (1)	33mmxxddnn	3mxdn
163	nn4wx	nnnn44wvxx	n4wx
164	c353e	cc335533cc	c353c
166	wddcp	wwdddccpp	wdcp
167	w52fn	ww555fffnn	w5fn
173	x775w	xx777755ww	x75w
174	e667x	ee666677xx	e67x
178	bxxfc	bbxxxxffcc	bxfc
181	gymn	ggyymmmnnn	gymn
184	my5n	mmmmyy55nn	my5n
191	8nn73	88nnnn7733	8n73
194	wgnwp	wwggmmwvpp	wgmwp
200	3nnpw	33nnnnppww	3npw
201	np4e	nnnnpp44ee	np4e
206	ncw7	nnccwww77	ncw7
207	ddxpp	dddxxpppp	dxp
208	g2577	gg22557777	g257
211	x7746	xx77774466	x746
220	e5n66	ee55nn6666	e5n6
226	44xe8	4444xxee88	4xe8
227	dnne7	ddnnnnnee77	dne7
231	ny5e	nnnnyy55ee	ny5e
237	gy433	ggyy443333	gy43
238	d66cn	dd6666ccnn	d6cn
239	nn6w6	nnnn66ww66	n6w6
244	pmg55	ppmmgg5555	pmg5
254	wmn6	wwwmmnn66	wmn6
256	d666m	dd6----6mm	d66m
262	nxx25	nnxxxx2255	nx25

```
[50]: print(mistakes_df['prediction_corrected'].str.len().value_counts())
```

```
prediction_corrected
```

```

4    39
3     6
5     5
Name: count, dtype: int64

```


```
[51]: mask = mistakes_df['prediction_corrected'].str.len() == 5
      mistakes_df[mask]
```

```
[51]:
```

	actual	prediction	prediction_corrected
62	cdcb3	eeddcbb33	edcb3
144	3n7mx (1)	33nn77mmxx	3n7mx
147	3mxdn (1)	33mmxxddnn	3mxdn
164	c353e	cc335533cc	c353c
194	wgmwp	wggmmwwpp	wgmwp

```
[52]: mistake_image_fp = os.path.join(data_path, mistakes_df[mask]['actual'].
      ↪values[0] + ".png")
      print(mistake_image_fp)
      mistake_image = Image.open(mistake_image_fp)
      plt.imshow(mistake_image)
      plt.show()
```

/content/drive/MyDrive/archive (1)/samples/cdcb3.png



```
[53]: train_accuracy = accuracy_score(results_train['actual'],
      ↪results_train['prediction_corrected'])
      print(train_accuracy)
      test_accuracy = accuracy_score(results_test['actual'],
      ↪results_test['prediction_corrected'])
      print(test_accuracy)
```

```

0.8580246913580247
0.8148148148148148

```